


# Tennis – Collaboration and Competition

Using Deep reinforcement learning for training multiple agent to compete.

Mohamed Dawod

dawod.moh@gmail.com



## Context:

In this project we train two agents to play tennis by controlling a racket to bounce a ball over net.

The environment state and reward scheme

a- Reward:

The agent is reward +0.1 for hitting the ball over the net, and -0.01 if the agent let the ball hit the ground or hit the ball out of bounds.

Thus, the goal for each agent is to keep the ball in the play.

b- State:

The observation space consists of 8 variables corresponding to the position and velocity of the ball and racket. Each agent receives its own, local observation. Two continuous actions are available, corresponding to movement toward (or away from) the net, and jumping.

c- Score:

This task is episodic, and the score for each episode is calculated by taking the maximum score of both agents. The agent is considered successfully trained if it managed to get an average score of +0.5 (over 100 consecutive episodes).

## The Solution:

To solve this project, we used the method of Multi Agent Deep Deterministic Policy Gradient (MADDPG) to solve this project, where each agent player has its own actor and critic network. However both agents are sharing the same random replay buffer. Sharing the same replay buffer, increase the speed of learning as it collects a massive amount of experiences really quickly.

## Deep Neural Network Model

The algorithm uses two deep neural networks (actor-critic) with the following structure:

- Actor
  - Hidden: (State\_size, 256) - ReLU
  - Hidden: (256, 128) - ReLU
  - Output: (128, 4) - TanH
- Critic
  - Hidden: (State\_size, 256) - ReLU
  - Hidden: (256 + action\_size, 128) - ReLU
  - Output: (128, 1) - Linear


## Hyperparameters

<code>BUFFER_SIZE = int(1e6)</code>	<code># replay buffer size</code>
<code>BATCH_SIZE = 512</code>	<code># minibatch size</code>
<code>GAMMA = 0.99</code>	<code># discount factor</code>
<code>TAU = 1e-3</code>	<code># for soft update of target parameters</code>
<code>LR_ACTOR = 1e-4</code>	<code># learning rate of the actor</code>
<code>LR_CRITIC = 3e-4</code>	<code># learning rate of the critic</code>
<code>WEIGHT_DECAY = 0</code>	<code># L2 weight decay</code>
<code>ACTOR_HL_SIZE = [256, 128]</code>	<code># Actor Hidden layers</code>
<code>CRITIC_HL_SIZE = [256, 128]</code>	<code># Critic Hidden layers</code>
<code>n_episodes = 3000, max_t = 2000</code>	<code># Number of episodes and times steps</code>

## Results

Episode 100	Current Score: 0.00	Average Score: 0.01
Episode 200	Current Score: 0.10	Average Score: 0.03
Episode 300	Current Score: 0.09	Average Score: 0.03
Episode 400	Current Score: 0.00	Average Score: 0.04
Episode 500	Current Score: 0.00	Average Score: 0.05
Episode 600	Current Score: 0.00	Average Score: 0.04
Episode 700	Current Score: 0.09	Average Score: 0.07
Episode 800	Current Score: 0.00	Average Score: 0.10
Episode 900	Current Score: 0.10	Average Score: 0.11
Episode 1000	Current Score: 0.10	Average Score: 0.18
Episode 1100	Current Score: 0.20	Average Score: 0.36
Episode 1200	Current Score: 0.30	Average Score: 0.38
Episode 1261	Current Score: 2.10	Average Score: 0.51

Environment solved in 1261 episodes! Average Score: 0.51  
Environment solved in 3182.52 seconds



## Future Work

I'm willing to try to improve the algorithm by implementing the following concepts:

- 1- [Prioritized Experience Replay](#)
- 2- [Exploration with Parameter Noise](#)